THE ONTOLOGY IMPERATIVE

Trustworthy Agentic AI for Boards and Executive Teams

Board Briefing

Agentic AI Liability and the Executive Mandate

Frédéric Verhelst · Version 2.0 · 15 June 2026

Your AI agents no longer just answer questions; they take actions, moving money, changing settings, sending messages, on their own and faster than anyone can review, and the organisation is fully liable for what they do. Accuracy is not the issue: these systems work by prediction, not by fixed rules, so even an agent working from perfect data can act outside what it was ever allowed to do. The controls most companies rely on were built for people working at human pace; what is missing is a fourth line that fires before an agent acts and can halt it, held by an executive senior enough to stop a launch the business is pushing.

TRADITIONAL GOVERNANCE *Three lines of defence*	REQUIRED ADDITION *The fourth line*
① **Operational management** ② **Risk and compliance** ③ **Internal audit**	④ **Structural enforcement** Constraints fire before action. Audit trail records every decision. **CDO custody.**
Human tempo · Reviews AFTER action	**Machine speed · Fires before, halts during**

That fourth line cannot be enforced from a subordinate position. The board charters it to a Chief Data Officer reporting to the CEO, the only seat that can hold a deployment gate against an executive pushing to go live. After any agent-driven loss, the first question a board faces is whether it governed a known risk. These six reveal whether the line exists today.

#	Six questions for the executive team
Q1	Who is chartered to block an unsafe agent at the deployment gate, before it reaches production?
Q2	Once an agent is live, who can halt it within seconds, without convening a committee?
Q3	Can we update agent constraints on our schedule, not the vendor's release cycle?
Q4	Can we reconstruct what an agent was authorised to do within 24 hours of an incident?
Q5	Do we have independent evidence agents ran as specified, or only the vendor's word?
Q6	Who is accountable when an agent produces a consequential error?

An organisation that cannot answer them is running autonomous systems without a calculable damage envelope. **The liability is absolute; the mandate must match it.**

About the author

Frédéric Verhelst, PhD (Applied Physics, TU Delft), works with boards and executive teams on the governance of trustworthy agentic AI, across twenty-five years in data, AI, and regulated industrial operations, including the enterprise data mandate at TotalEnergies EP Danmark and, most recently, agentic AI in safety-critical maritime operations. He is available for non-executive director, board-advisory, and CDO appointments focused on corporate AI governance.

For the full picture, see the Governing Agentic AI board pack and the Agentic AI Capability Stack™, at fredericverhelst.com/TOI-library. Licensed CC BY-ND 4.0 with attribution.